

## Lab Guide 9

### MPI-Based Parallelization

#### Objective:

- Learn the basics of program parallelization with distributed memory programming (message passing)
- Understand the basic types of parallelization: 1) functional versus 2) data parallelization

#### Introduction

This lab session aims to apply the basic MPI communication concepts studied in the previous session to parallelize simple applications.

Copy the `/share/cpar/P09_Codigo` folder to your home directory in the SeARCH cluster.

Compile the program in the cluster frontend using the `mpic++ -O2 -o primes PrimeMain.cpp` command. Use the `sbatch primes_mpi.sh` command to run the application.

The `primes_mpi.sh` file should specify the required resources and should run the MPI application.

The following example requests three PUs and spawns three MPI processes:

```
[search7edu]$ cat primes_mpi.sh
#!/bin/bash
#SBATCH --time=1:00
#SBATCH --ntasks=3
#SBATCH --partition=cpar
mpirun -np 3 ./primes
```

The number of requested resources (`--ntasks`) must be the same as the number of processes (`-np`) used in the `mpirun` command.

#### Exercise 1 – Prime calculation using the Sieve of Eratosthenes

Consider the following sequential program, which finds all prime numbers up to a given `MAXP`:

```
int MAXP = 1000000;
int SMAXP = 1000;
int pack=MAXP/10;

PrimeServer *ps1 = new PrimeServer();
PrimeServer *ps2 = new PrimeServer();
PrimeServer *ps3 = new PrimeServer();

ps1->minitFilter(1, SMAXP/3, SMAXP);
ps2->minitFilter(SMAXP/3+1, 2*SMAXP/3, SMAXP);
ps3->minitFilter(2*SMAXP/3+1, SMAXP, SMAXP);

int *ar = new int[pack/2];
for(int i=0; i<10; i++) {
    generate(i*pack, (i+1)*pack, ar);
    ps1->mprocess(ar, pack/2);
    ps2->mprocess(ar, pack/2);
    ps3->mprocess(ar, pack/2);
}
ps3->end();
```


- Parallelize the code using MPI through the implementation of a pipeline of processes that receives an array of integers, created by the `generate` function, and each process filters out a subset of the input. The `mprocess` method implements the filtering of the primes, and `end` prints the final amount of primes found. This pipeline should have 3 processes, one for each instance of `PrimeServer` performing the filtering.
- (\*) Modify the parallelization implemented in **a)** to work with an arbitrary number of processes and messages.
- (\*) Parallelize the sequential application through the implementation of a farm of processes behaving in a “work sharing” paradigm with dynamic scheduling.

## Exercise 2 – Distributed memory matrix multiplication with MPI collective operations

Download the matrix multiplication code from lab session 0 (base DOT version in the figure).

```
foreach line of A
  foreach column of B
    C[line,column] = DOT[line of A,column of B]
```

$$C_{ij} = DOT_{linha\_A_i,coluna\_B_j} = \sum_{k=0}^{n-1} (A_{ik} * B_{kj})$$



- Identify a strategy to implement the matrix multiplication with MPI:
  - Each process will compute a subset of the elements of matrix C. Identify the main alternatives to partition the computation of this matrix among processes.
  - Select the simplest alternative to partition the computations of C. What data from matrix A and matrix B is required to perform the computation of each part of C?
- Implement the selected approach using MPI collective operations. The relevant collective operations are:

- Broadcast** the data from root to all other processes:

```
int MPI_Bcast(void* buffer, int count, MPI_Datatype
datatype, int root, MPI_Comm comm)
```


- Scatter**: scatters data from root into all other processes

```
int MPI_Scatter(void* sbuf, int scount, MPI_Datatype
stype, void* rbuf, int rcount, MPI_Datatype rtype, int
root, MPI_Comm comm )
```


- Gather**: Joins data from all processes into the root

```
int MPI_Gather(/* same signature of Scatter*/)

```



MPI\_Bcast



MPI\_Scatter

