



University of Minho
School of Engineering



Dados e Aprendizagem Automática

Artificial Neural Networks:

Multilayer Perceptron

DAA @ MEI-1º/MiEI-4º – 1º Semestre

Bruno Fernandes, César Analide, Dalila Alves, Filipa Ferraz, Victor Alves

Part VIII


Contents

2

- Artificial Neural Networks
 - Multilayer Perceptron
- Hands On

3

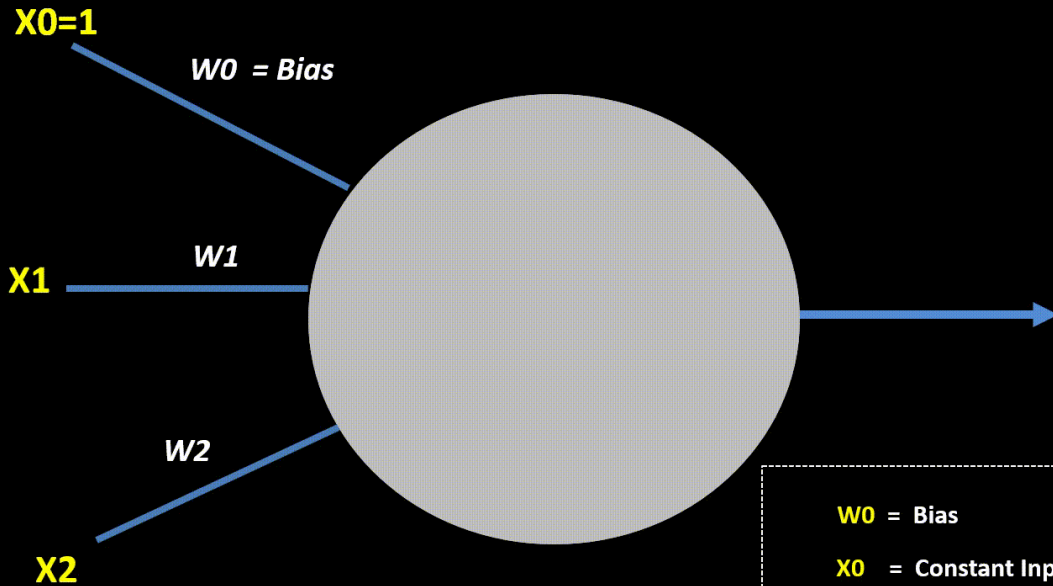
Artificial Neural Networks



Artificial Neural Networks

5


Artificial Neuron



$W0 = \text{Bias}$
 $X0 = \text{Constant Input 1 for Bias}$
 $X1, X2 = \text{Attribute Inputs}$
 $W1, W2 = \text{Weights of Inputs } X1, X2$

machinelearningknowledge.ai

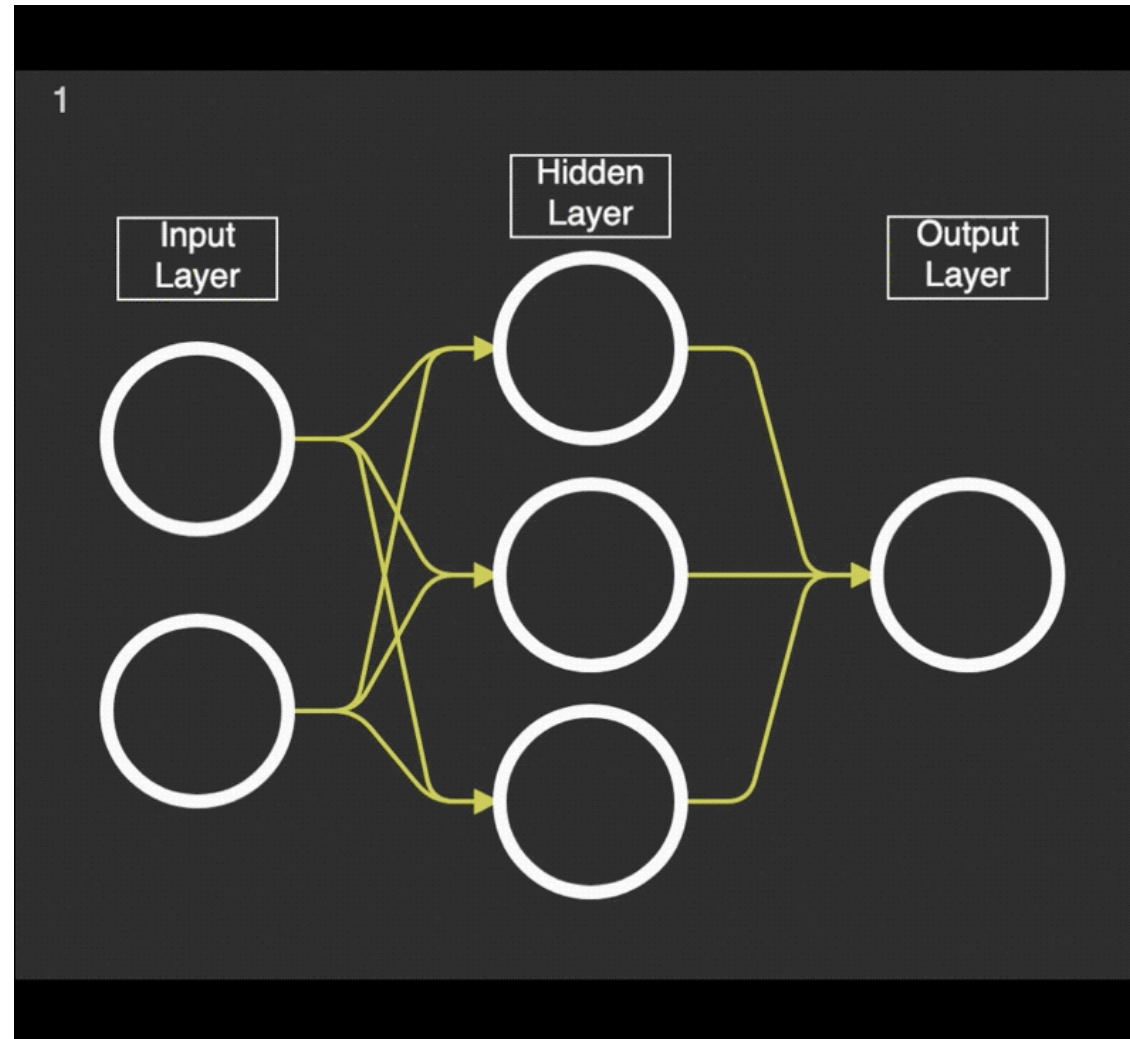
Neural Network – Backpropagation



© machinelearningknowledge.ai

Artificial Neural Networks

6



Practical Application of Artificial Neural Networks

7

For this example, we will use the already known dataset about USA housing. This is the real estate problem, and our goal is to help the real estate agent predict housing prices for regions in the USA (dataset [here](#)).

We have used Linear Regression in this context; now, we are going to try **Artificial Neural Networks**. Let's try using **Multilayer Perceptrons (MLPs)**, a class of **Artificial Neural Networks**!

To implement our first **Artificial Neural Network**, we will use:

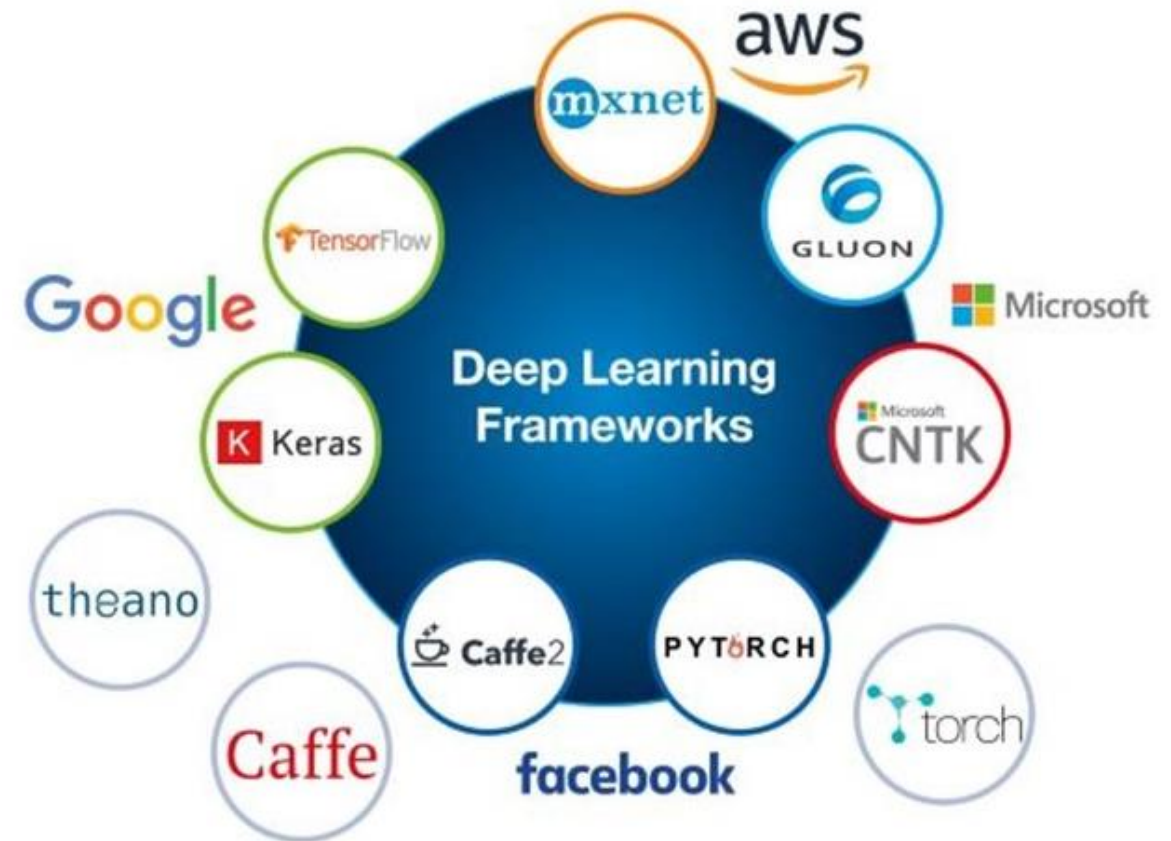


An open source machine learning
library for research and
production.

Remembering Frameworks

8

Theano	Python library that allows you to define, optimize and calculate mathematical expressions with multidimensional arrays efficiently. It works with GPUs and efficiently performs differential calculations. <i>University of Montreal's lab, MLA</i>
Lasagne	Light library for building and training neural networks using Theano
Blocks	Theano-based framework for building and training neural networks
TensorFlow	Open-source library for numerical computation using graphs <i>Google Brain team</i>
Keras	Deep learning library for Python. Runs on Theano or TensorFlow
MXNet	Deep learning framework designed for efficiency and flexibility <i>Amazon</i>
PyTorch	Flexible tensors and neural networks with strong GPU support <i>Facebook Artificial Intelligence Research team (FAIR)</i>
Torch	<i>Ronan Collobert</i>
Caffe	<i>Berkeley Vision and Learning Center</i>
CNTK	<i>Microsoft</i>
Deeplearning4j	<i>SkyMind</i>

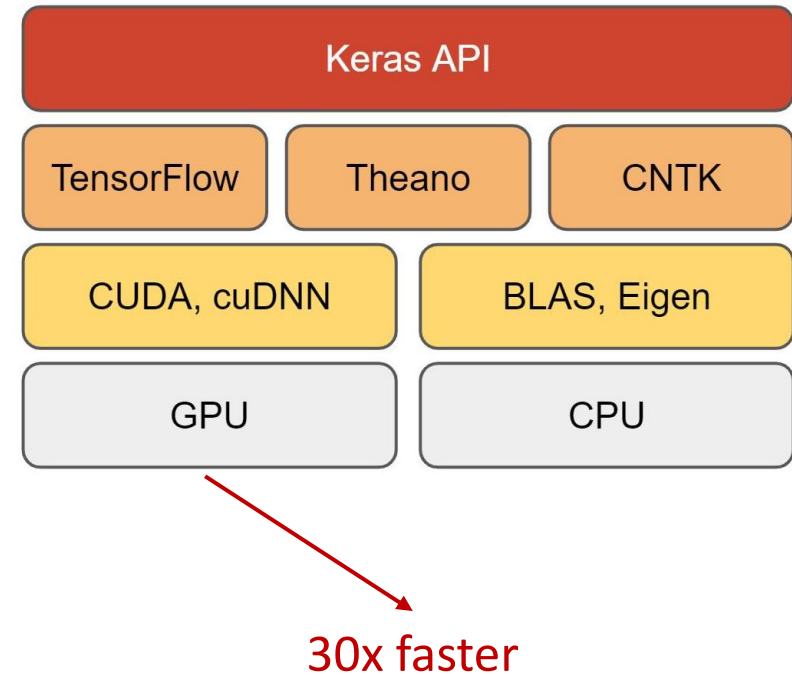


TensorFlow

9

Why?

- Open-source software library for **high-performance numerical computation**
- Strong support for **machine learning** and **deep learning**
- It has seen tremendous growth and popularity in the machine learning community



Implementing a Multilayer Perceptron

10

The data

It will be used data frame with 5000 observations on the following 7 variables:

- **Avg. Area Income** - Avg. Income of residents of the city house is located in.
- **Avg. Area House Age** - Avg Age of Houses in same city
- **Avg. Area Number of Rooms** - Avg Number of Rooms for Houses in same city
- **Avg. Area Number of Bedrooms** - Avg Number of Bedrooms for Houses in same city
- **Area Population** - Population of city house is located in
- **Price** - Price that the house sold at
- **Address** - Address for the house

Implementing a Multilayer Perceptron

11

Import libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

Get the data

Create the data frame

```
USAhousing = pd.read_csv('USA_Housing.csv')
```

```
USAhousing.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 5000 entries, 0 to 4999
```

```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	Avg. Area Income	5000 non-null	float64
1	Avg. Area House Age	5000 non-null	float64
2	Avg. Area Number of Rooms	5000 non-null	float64
3	Avg. Area Number of Bedrooms	5000 non-null	float64
4	Area Population	5000 non-null	float64
5	Price	5000 non-null	float64
6	Address	5000 non-null	object

```
dtypes: float64(6), object(1)
```

```
memory usage: 273.6+ KB
```

Implementing a Multilayer Perceptron

12

Since the feature *Address* is the one categoric and not needed for the purpose of the exercise, let's drop it:

```
USAhousing.drop('Address', axis = 1, inplace = True)
```

```
USAhousing.head()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
0	79545.458574	5.682861	7.009188	4.09	23086.800503	1.059034e+06
1	79248.642455	6.002900	6.730821	3.09	40173.072174	1.505891e+06
2	61287.067179	5.865890	8.512727	5.13	36882.159400	1.058988e+06
3	63345.240046	7.188236	5.586729	3.26	34310.242831	1.260617e+06
4	59982.197226	5.040555	7.839388	4.23	26354.109472	6.309435e+05

Implementing a Multilayer Perceptron

13

```
USAhousing.describe()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
50%	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
75%	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
max	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06

Implementing a Multilayer Perceptron


14

EDA

Create the histogram of the target - column *Price*:

```
sns.histplot(USAhousing['Price'])
```

<Axes: xlabel='Price', ylabel='Count'>



Create a heatmap of the features:



Create a pairplot to visualize relations:

Implementing a Multilayer Perceptron

15

Train Test Split

Define X and Y :

```
X = USAhousing.drop('Price', axis = 1)
y = USAhousing[['Price']]
```

Divide the subsets of test and training data:

```
from sklearn.model_selection import GridSearchCV, KFold, train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 2021)
```

Implementing a Multilayer Perceptron

16

To implement **our first MLP** we will need some more libraries! Let's import them all at once! You'll need to install **TensorFlow**, **Keras** and **Scikeras** - use the Navigator or the Prompt:

```
conda install -c conda-forge tensorflow
conda install -c conda-forge keras
conda install -c conda-forge scikeras
```

or
or
or

```
pip install tensorflow
pip install --upgrade keras
pip install scikeras[tensorflow]
```

Import more libraries

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from scikeras.wrappers import KerasRegressor
```

```
print("TensorFlow version:", tf.__version__)
```

TensorFlow version: 2.14.0

Implementing a Multilayer Perceptron

17

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

Structure the MLP

Define a model with:

- *ReLU* as activation function
- sequential topology
- three layers
- *MAE* as loss function
- *Adam* as optimizer
- learning rate of *0.01*
- *MAE* and *MSE* as metrics

```
def build_model(activation = 'relu', learning_rate = 0.01):  
    model = Sequential()  
    model.add(Dense(16, input_dim = 5, activation = activation))  
    model.add(Dense(8, activation = activation))  
    model.add(Dense(1, activation = activation)) #output  
  
    #Compile the model  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Implementing a Multilayer Perceptron

18


Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

Structure the MLP

Define a model with:

- *ReLU* as activation function
- sequential topology
- three layers
- *MAE* as loss function
- *Adam* as optimizer
- learning rate of *0.01*
- *MAE* and *MSE* as metrics

```
def build_model(activation = 'relu', learning_rate = 0.01):  
    model = Sequential()  
    model.add(Dense(16, input_dim = 5, activation = activation))  
    model.add(Dense(8, activation = activation))  
    model.add(Dense(1, activation = activation)) #output  
  
    #Compile the model  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```



Layers stacked one
over the other

Implementing a Multilayer Perceptron

19

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

Structure the MLP

Define a model with:

- *ReLU* as activation function
- sequential topology
- three layers
- *MAE* as loss function
- *Adam* as optimizer
- learning rate of *0.01*
- *MAE* and *MSE* as metrics

```
def build_model(activation = 'relu', learning_rate = 0.01):  
    model = Sequential()  
    model.add(Dense(16, input_dim = 5, activation = activation))  
    model.add(Dense(8, activation = activation))  
    model.add(Dense(1, activation = activation)) #output  
  
    #Compile the model  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Layers stacked one
over the other

Number of neurons in
each layer
(last one is the output)

Implementing a Multilayer Perceptron

20

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

Structure the MLP

Define a model with:

- *ReLU* as activation function
- sequential topology
- three layers
- *MAE* as loss function
- *Adam* as optimizer
- learning rate of *0.01*
- *MAE* and *MSE* as metrics

```
def build_model(activation = 'relu', learning_rate = 0.01):  
    model = Sequential()  
    model.add(Dense(16, input_dim = 5, activation = activation))  
    model.add(Dense(8, activation = activation))  
    model.add(Dense(1, activation = activation)) #output  
  
    #Compile the model  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Layers stacked one
over the other

Number of neurons in
each layer
(last one is the output)

Number of input
features

Implementing a Multilayer Perceptron

21

Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

Structure the MLP

Define a model with:

- *ReLU* as activation function
- sequential topology
- three layers
- *MAE* as loss function
- *Adam* as optimizer
- learning rate of *0.01*
- *MAE* and *MSE* as metrics

```
def build_model(activation = 'relu', learning_rate = 0.01):  
    model = Sequential()  
    model.add(Dense(16, input_dim = 5, activation = activation))  
    model.add(Dense(8, activation = activation))  
    model.add(Dense(1, activation = activation)) #output  
  
    #Compile the model  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```

Layers stacked one
over the other

Number of neurons in
each layer
(last one is the output)

Number of input
features

Activation function
(here as an argument to
the *build_model* function).

Implementing a Multilayer Perceptron

22

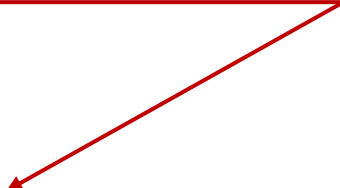
Let's build our model using the **Sequential API** (there are others but, for now, let's keep it simple)!

Structure the MLP

Define a model with:

- *ReLU* as activation function
- sequential topology
- three layers
- *MAE* as loss function
- *Adam* as optimizer
- learning rate of *0.01*
- *MAE* and *MSE* as metrics

```
def build_model(activation = 'relu', learning_rate = 0.01):  
    model = Sequential()  
    model.add(Dense(16, input_dim = 5, activation = activation))  
    model.add(Dense(8, activation = activation))  
    model.add(Dense(1, activation = activation)) #output  
  
    #Compile the model  
    model.compile(  
        loss = 'mae',  
        optimizer = tf.optimizers.Adam(learning_rate),  
        metrics = ['mae', 'mse'])  
    return model
```



After building the stacked MLP, we need to **compile the model** by setting the **loss function** (MAE as we are solving a regression problem), the **optimizer** (which implements the gradient descent and updates the weights), and a set of **metrics** (to further understand the performance of the model - not used when backpropagating the error).

Implementing a Multilayer Perceptron

23

Build the model:

```
model = build_model()  
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
dense (Dense)	(None, 16)	96
dense_1 (Dense)	(None, 8)	136
dense_2 (Dense)	(None, 1)	9
=====		
Total params: 241 (964.00 Byte)		
Trainable params: 241 (964.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		

Implementing a Multilayer Perceptron

24

GridSearchCV

We want to **find the best possible MLP** to solve our problem so... Let's **tune it** (at least, some hyperparameters)!

We must first define a dictionary of {key -> list of values}. For now, we will only **tune** the **optimizer** (trying three values - so, **we will fit 3 candidate models**). Note that these are arguments of the *build_model()* function.

Define the grid parameters using a dictionary:

```
optimizer = ['SGD', 'RMSprop', 'Adagrad']  
param_grid = dict(optimizer = optimizer)
```

Define a **KFold with 5 splits, shuffle and a random state:**

```
kf = KFold(n_splits = 5, shuffle = True, random_state = 2021)
```

Use a *KerasRegressor* with a *batch size* of 32, *validation split* of 0.2 and 20 *epochs*:

```
model = KerasRegressor(model = build_model, batch_size = 32, validation_split = 0.2, epochs = 20)
```

Compute a *GridSearchCV* with *NegMAE scoring*, *refit* and a *verbose* of 1:

```
grid_search = GridSearchCV(estimator = model, param_grid = param_grid, cv = kf, scoring = 'neg_mean_absolute_error', refit = 'True', verbose = 1)
```


Implementing a Multilayer Perceptron

25

GridSearchCV

We want to **find the best possible MLP** to solve our problem so... Let's **tune it** (at least, some hyperparameters)!

We must first define a dictionary of {key -> list of values}. For now, we will only **tune** the **optimizer** (trying three values - so, **we will fit 3 candidate models**). Note that these are arguments of the `build_model()` function.

Define the grid parameters using a dictionary:

```
optimizer = ['SGD', 'RMSprop', 'Adagrad']  
param_grid = dict(optimizer = optimizer)
```

Define a *KFold* with 5 splits, shuffle and a random state:

```
kf = KFold(n_splits = 5, shuffle = True, random_state = 2021)
```

Use a **KerasRegressor** with a *batch size* of 32, *validation split* of 0.2 and 20 epochs:

```
model = KerasRegressor(model = build_model, batch_size = 32, validation_split = 0.2, epochs = 20)
```

Compute a *GridSearchCV* with *NegMAE* scoring, *refit* and a *verbose* of 1:

```
grid_search = GridSearchCV(estimator = model, param_grid = param_grid, cv = kf, scoring = 'neg_mean_absolute_error', refit = 'True', verbose = 1)
```

We must pass our **build_model** function, define the number of **epochs** (*number of passes of the entire training dataset*) and the **batch size** (*number of training samples in one forward/backward pass*). The KerasRegressor API will **return an instance of our MLP**.

Implementing a Multilayer Perceptron

26

GridSearchCV

We want to **find the best possible MLP** to solve our problem so... Let's **tune it** (at least, some hyperparameters)!

We must first define a dictionary of {key -> list of values}. For now, we will only **tune** the **optimizer** (trying three values - so, **we will fit 3 candidate models**). Note that these are arguments of the `build_model()` function.

Define the grid parameters using a dictionary:

```
optimizer = ['SGD', 'RMSprop', 'Adagrad']  
param_grid = dict(optimizer = optimizer)
```

Define a *KFold* with 5 splits, shuffle and a random state:

```
kf = KFold(n_splits = 5, shuffle = True, random_state = 2021)
```


Use a *KerasRegressor* with a batch size of 32, validation split of 0.2 and 20 epochs:

```
model = KerasRegressor(model = build_model, batch_size = 32, validation_split = 0.2, epochs = 20)
```

Compute a *GridSearchCV* with *NegMAE* scoring, refit and a verbose of 1:

```
grid_search = GridSearchCV(estimator = model, param_grid = param_grid, cv = kf, scoring = 'neg_mean_absolute_error', refit = 'True', verbose = 1)
```

Fraction of the training data to be used as **validation data**. The model will set apart this fraction of the training data, **will not train on it**, and **will evaluate the loss and the model metrics** on this data at the end of each epoch.



Implementing a Multilayer Perceptron

27

GridSearchCV

We want to **find the best possible MLP** to solve our problem so... Let's **tune it** (at least, some hyperparameters)!

We must first define a dictionary of *{key -> list of values}*. For now, we will only **tune** the **optimizer** (trying three values - so, **we will fit 3 candidate models**). Note that these are arguments of the *build_model()* function.

Define the grid parameters using a dictionary:

```
optimizer = ['SGD', 'RMSprop', 'Adagrad']  
param_grid = dict(optimizer = optimizer)
```

Define a *KFold* with 5 splits, shuffle and a random state:

```
kf = KFold(n_splits = 5, shuffle = True, random_state = 2021)
```

Use a *KerasRegressor* with a batch size of 32, validation split of 0.2 and 20 epochs:

```
model = KerasRegressor(model = build_model, batch_size = 32, validation_split = 0.2, epochs = 20)
```

Compute a GridSearchCV with NegMAE scoring, refit and a verbose of 1:

```
grid_search = GridSearchCV(estimator = model, param_grid = param_grid, cv = kf, scoring = 'neg_mean_absolute_error', refit = 'True', verbose = 1)
```

Implementing a Multilayer Perceptron

28

Fit the model:

```
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 3 candidates, totalling 15 fits

Epoch 1/20

```
80/80 [=====] - 2s 8ms/step - loss: 1227161.2500 - mae: 1227161.2500 - mse: 1634915057664.0000  
- val_loss: 1247080.1250 - val_mae: 1247080.1250 - val_mse: 1679838281728.0000
```

Epoch 2/20

```
80/80 [=====] - 0s 4ms/step - loss: 1227161.1250 - mae: 1227161.1250 - mse: 1634914664448.0000  
- val_loss: 1247080.1250 - val_mae: 1247080.1250 - val_mse: 1679838281728.0000
```

Epoch 3/20

```
80/80 [=====] - 0s 4ms/step - loss: 1227161.1250 - mae: 1227161.1250 - mse: 1634914795520.0000  
- val_loss: 1247080.1250 - val_mae: 1247080.1250 - val_mse: 1679838281728.0000
```

Epoch 4/20

```
80/80 [=====] - 0s 3ms/step - loss: 1227161.3750 - mae: 1227161.3750 - mse: 1634914795520.0000  
- val_loss: 1247080.1250 - val_mae: 1247080.1250 - val_mse: 1679838281728.0000
```

Epoch 5/20

```
80/80 [=====] - 0s 3ms/step - loss: 1227161.2500 - mae: 1227161.2500 - mse: 1634914402304.0000  
- val_loss: 1247080.1250 - val_mae: 1247080.1250 - val_mse: 1679838281728.0000
```

Epoch 6/20

```
80/80 [=====] - 0s 3ms/step - loss: 1227161.2500 - mae: 1227161.2500 - mse: 1634914795520.0000  
- val_loss: 1247080.1250 - val_mae: 1247080.1250 - val_mse: 1679838281728.0000
```

Epoch 7/20

```
80/80 [=====] - 0s 4ms/step - loss: 1227161.5000 - mae: 1227161.5000 - mse: 1634914533376.0000  
- val_loss: 1247080.1250 - val_mae: 1247080.1250 - val_mse: 1679838281728.0000
```

Epoch 8/20

...

Implementing a Multilayer Perceptron

29

We can now **analyze the performance** of our MLP:

```
print("Best: %f using %s" % (grid_search.best_score_, grid_search.best_params_))
```

```
Best: -202760.793380 using {'optimizer': 'Adagrad'}
```

→ The best model

Find the *mean test score*, *std test score* and *params* for each search:

```
means = grid_search.cv_results_['mean_test_score']  
stds = grid_search.cv_results_['std_test_score']  
params = grid_search.cv_results_['params']
```

```
for mean, stdev, param in zip(means, stds, params):  
    print("%f (%f) with: %r" % (mean, stdev, param))
```

```
-614376.884116 (499562.792066) with: {'optimizer': 'SGD'}  
-408853.059894 (411606.908481) with: {'optimizer': 'RMSprop'}  
-202760.793380 (1187.979771) with: {'optimizer': 'Adagrad'}
```

Implementing a Multilayer Perceptron

30

```
best_mlp_model = grid_search.best_estimator_  
print(best_mlp_model)  
  
KerasRegressor(  
    model=<function build_model at 0x000002E011306660>  
    build_fn=None  
    warm_start=False  
    random_state=None  
    optimizer=Adagrad  
    loss=None  
    metrics=None  
    batch_size=32  
    validation_batch_size=None  
    verbose=1  
    callbacks=None  
    validation_split=0.2  
    shuffle=True  
    run_eagerly=False  
    epochs=20  
)
```

Implementing a Multilayer Perceptron

31

Let's try to **understand if our model overfitted**:

```
best_mlp_model.fit(X_train, y_train, epochs = 20, validation_data = (X_test, y_test), verbose = 1)
```

Epoch 1/20

```
125/125 [=====] - 2s 5ms/step - loss: 561256.3125 - mae: 561256.3125 - mse: 559574679552.0000  
0 - val_loss: 195579.8438 - val_mae: 195579.8438 - val_mse: 59914842112.0000
```

Epoch 2/20

```
125/125 [=====] - 0s 3ms/step - loss: 203607.3906 - mae: 203607.3906 - mse: 64345059328.0000  
- val_loss: 198204.0000 - val_mae: 198204.0000 - val_mse: 61535883264.0000
```

Epoch 3/20

```
125/125 [=====] - 0s 3ms/step - loss: 203205.4844 - mae: 203205.4844 - mse: 63991103488.0000  
- val_loss: 195700.2344 - val_mae: 195700.2344 - val_mse: 60022448128.0000
```

Epoch 4/20

```
125/125 [=====] - 0s 3ms/step - loss: 202106.3750 - mae: 202106.3750 - mse: 63488376832.0000  
- val_loss: 198870.7969 - val_mae: 198870.7969 - val_mse: 61997703168.0000
```

Epoch 5/20

```
125/125 [=====] - 0s 3ms/step - loss: 204026.1250 - mae: 204026.1250 - mse: 64430120960.0000  
- val_loss: 196261.9844 - val_mae: 196261.9844 - val_mse: 60319195136.0000
```

Epoch 6/20

```
125/125 [=====] - 0s 3ms/step - loss: 203777.8438 - mae: 203777.8438 - mse: 64410632192.0000  
- val_loss: 195595.4062 - val_mae: 195595.4062 - val_mse: 59969236992.0000
```

Epoch 7/20

```
125/125 [=====] - 0s 4ms/step - loss: 203254.9219 - mae: 203254.9219 - mse: 64242946048.0000  
- val_loss: 198508.5781 - val_mae: 198508.5781 - val_mse: 61783752704.0000
```

Epoch 8/20

```
125/125 [=====] - 0s 3ms/step - loss: 203353.2812 - mae: 203353.2812 - mse: 63824674816.0000  
- val_loss: 197179.8750 - val_mae: 197179.8750 - val_mse: 60830883840.0000
```

...

Implementing a Multilayer Perceptron

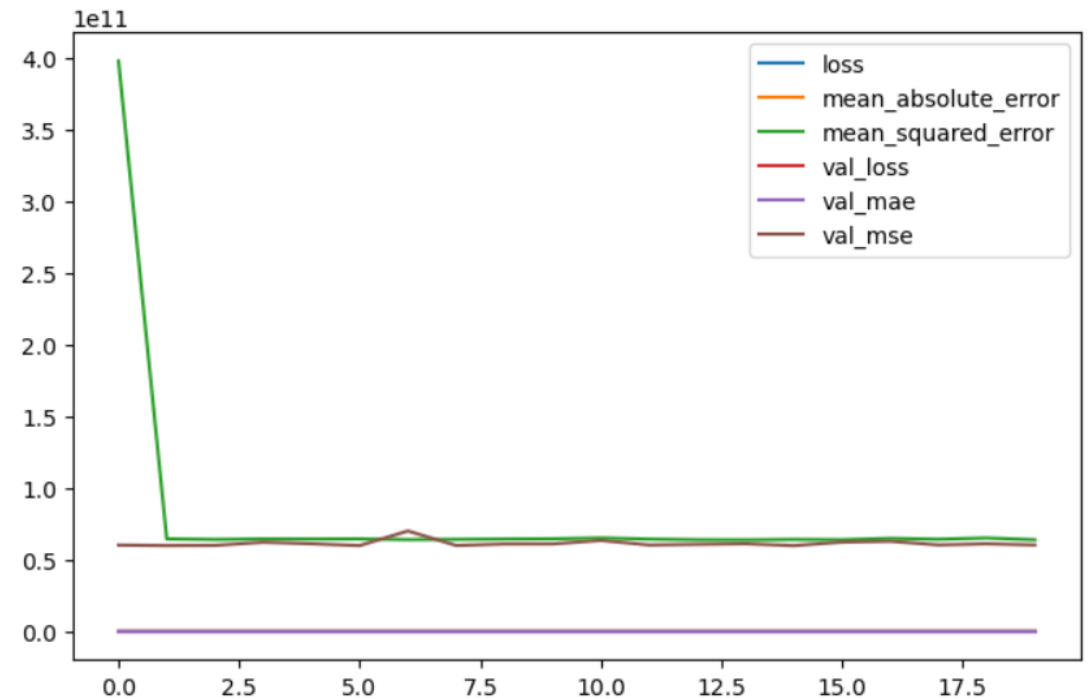

32

Let's plot it:

```
plt.plot(best_mlp_model.history_['loss'])
plt.plot(best_mlp_model.history_['val_loss'])
plt.title('model performance')
plt.ylabel('loss values')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```

Did the **model overfit**?

```
pd.DataFrame(best_mlp_model.history_).plot(figsize = (8, 5))
plt.show()
```



Implementing a Multilayer Perceptron

33

Predictions

Obtain the predictions:

```
predictions = best_mlp_model.predict(X_test)
```

```
32/32 [=====] - 0s 2ms/step
```

Print the first five:

```
array([[1394124. ],  
       [1224608.1],  
       [1049303.8],  
       [1335269.2],  
       [1235241. ]], dtype=float32)
```

Evaluate the model

```
from sklearn import metrics
```

```
print('MAE:', metrics.mean_absolute_error(y_test, predictions))  
print('MSE:', metrics.mean_squared_error(y_test, predictions))  
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test, predictions)))
```


MAE: 197376.8433273626

MSE: 61088757946.53201

RMSE: 247161.40059995616

Scatter the real values with the predictions:

```
plt.scatter(y_test, predictions)
```



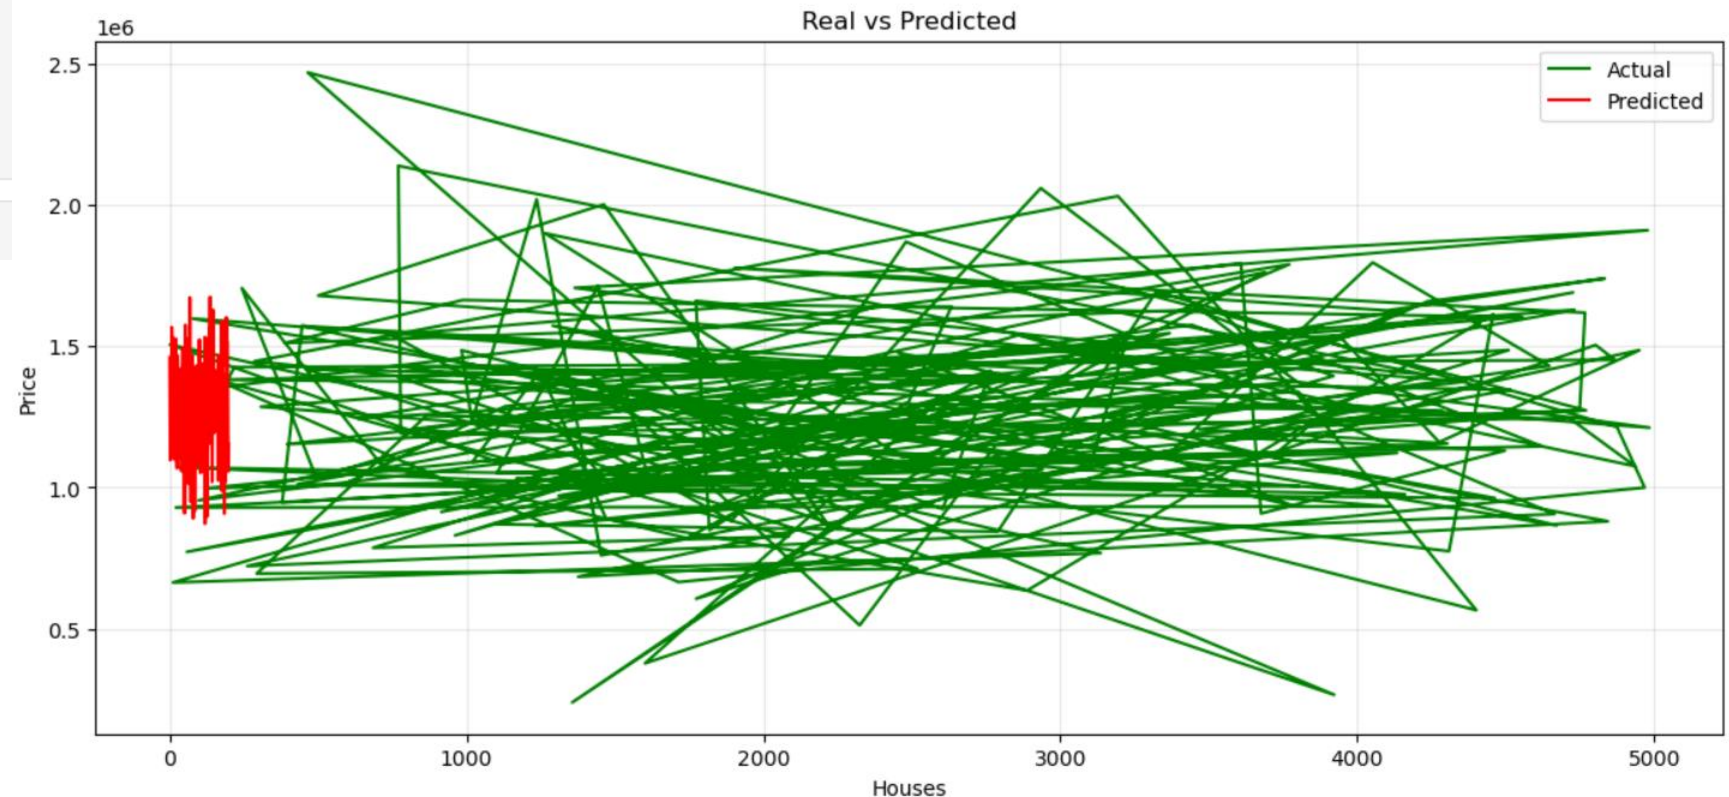
Implementing a Multilayer Perceptron

34

Create a visualization of the actual and predicted results. Limit it to 200 comparisons:

```
def real_predicted_viz(limit):  
    plt.figure(figsize = (14, 6))  
    plt.plot(y_test[:limit], color = 'green', label = 'Actual')  
    plt.plot(predictions[:limit], color = 'red', label = 'Predicted')  
    plt.grid(alpha = 0.3)  
    plt.xlabel('Houses')  
    plt.ylabel('Price')  
    plt.title('Real vs Predicted')  
    plt.legend()  
    plt.show()
```

```
real_predicted_viz(200)
```



Implementing a Multilayer Perceptron

35

Data scaling

Data scaling or **normalization** is a process of making model data in a standard format so that the training is improved, accurate, and faster.

```
USAhousing.describe()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population	Price
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5.000000e+03
mean	68583.108984	5.977222	6.987792	3.981330	36163.516039	1.232073e+06
std	10657.991214	0.991456	1.005833	1.234137	9925.650114	3.531176e+05
min	17796.631190	2.644304	3.236194	2.000000	172.610686	1.593866e+04
25%	61480.562388	5.322283	6.299250	3.140000	29403.928702	9.975771e+05
50%	68804.286404	5.970429	7.002902	4.050000	36199.406689	1.232669e+06
75%	75783.338666	6.650808	7.665871	4.490000	42861.290769	1.471210e+06
max	107701.748378	9.519088	10.759588	6.500000	69621.713378	2.469066e+06

Implementing a Multilayer Perceptron

36

Artificial neural networks are "picky" - they prefer scaled data! Therefore, and since our data have a large variation of values, let's scale the data to be in the interval between [0, 1]:

```
from sklearn.preprocessing import MinMaxScaler

scaler_X = MinMaxScaler(feature_range=(0, 1)).fit(X)
scaler_y = MinMaxScaler(feature_range=(0, 1)).fit(y)
X_scaled = pd.DataFrame(scaler_X.transform(X[X.columns]), columns = X.columns)
y_scaled = pd.DataFrame(scaler_y.transform(y[y.columns]), columns = y.columns)
```

Let's check the data **before and after** transformations.

Implementing a Multilayer Perceptron

37

```
X.head()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population
0	79545.458574	5.682861	7.009188	4.09	23086.800503
1	79248.642455	6.002900	6.730821	3.09	40173.072174
2	61287.067179	5.865890	8.512727	5.13	36882.159400
3	63345.240046	7.188236	5.586729	3.26	34310.242831
4	59982.197226	5.040555	7.839388	4.23	26354.109472

And now scaled:

```
X_scaled.head()
```

	Avg. Area Income	Avg. Area House Age	Avg. Area Number of Rooms	Avg. Area Number of Bedrooms	Area Population
0	0.686822	0.441986	0.501502	0.464444	0.329942
1	0.683521	0.488538	0.464501	0.242222	0.575968
2	0.483737	0.468609	0.701350	0.695556	0.528582
3	0.506630	0.660956	0.312430	0.280000	0.491549
4	0.469223	0.348556	0.611851	0.495556	0.376988

```
y.head()
```

	Price
0	1.059034e+06
1	1.505891e+06
2	1.058988e+06
3	1.260617e+06
4	6.309435e+05

And now scaled:

```
y_scaled.head()
```

	Price
0	0.425210
1	0.607369
2	0.425192
3	0.507384
4	0.250702

Implementing a Multilayer Perceptron

38

We need to **split the training and test** sets again:

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_scaled, test_size = 0.2, random_state = 2021)
```

And then recall the MLP:

```
model = build_model()
model.summary()
```

Model: "sequential_18"

Layer (type)	Output Shape	Param #
=====	=====	=====
dense_54 (Dense)	(None, 16)	96
dense_55 (Dense)	(None, 8)	136
dense_56 (Dense)	(None, 1)	9
=====	=====	=====
Total params: 241 (964.00 Byte)		
Trainable params: 241 (964.00 Byte)		
Non-trainable params: 0 (0.00 Byte)		
=====	=====	=====

```
model = KerasRegressor(model = build_model, batch_size = 32, validation_split = 0.2, epochs = 20)
```

```
grid_search = GridSearchCV(estimator = model, param_grid = param_grid, cv = kf, scoring = 'neg_mean_absolute_error', refit = 'True', verbose = 1)
```

Implementing a Multilayer Perceptron

39

Fit the model:

```
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 3 candidates, totalling 15 fits

Epoch 1/20

80/80 [=====] - 2s 7ms/step - loss: 0.0902 - mae: 0.0902 - mse: 0.0148 - val_loss: 0.0546 - val_mae: 0.0546
- val_mse: 0.0048

Epoch 2/20

80/80 [=====] - 0s 4ms/step - loss: 0.0416 - mae: 0.0416 - mse: 0.0028 - val_loss: 0.0344 - val_mae: 0.0344
- val_mse: 0.0018

Epoch 3/20

80/80 [=====] - 0s 4ms/step - loss: 0.0353 - mae: 0.0353 - mse: 0.0019 - val_loss: 0.0466 - val_mae: 0.0466
- val_mse: 0.0032

Epoch 4/20

80/80 [=====] - 0s 4ms/step - loss: 0.0363 - mae: 0.0363 - mse: 0.0021 - val_loss: 0.0352 - val_mae: 0.0352
- val_mse: 0.0019

Epoch 5/20

80/80 [=====] - 0s 4ms/step - loss: 0.0370 - mae: 0.0370 - mse: 0.0021 - val_loss: 0.0359 - val_mae: 0.0359
- val_mse: 0.0019

Epoch 6/20

80/80 [=====] - 0s 4ms/step - loss: 0.0391 - mae: 0.0391 - mse: 0.0024 - val_loss: 0.0345 - val_mae: 0.0345
- val_mse: 0.0018

Epoch 7/20

80/80 [=====] - 0s 4ms/step - loss: 0.0346 - mae: 0.0346 - mse: 0.0019 - val_loss: 0.0351 - val_mae: 0.0351
- val_mse: 0.0019

Epoch 8/20

...

Implementing a Multilayer Perceptron

40

Find the *best score* and the *best params*:

```
print("Best: %f using %s" % (grid_search.best_score_, grid_search.best_params_))  
Best: -0.219534 using {'optimizer': 'SGD'}
```

Find the *mean test score*, *std test score* and *params* for each search:

```
means = grid_search.cv_results_['mean_test_score']  
stds = grid_search.cv_results_['std_test_score']  
params = grid_search.cv_results_['params']  
  
for mean, stdev, param in zip(means, stds, params):  
    print("%f (%f) with: %r" % (mean, stdev, param))  
  
-0.219534 (0.225235) with: {'optimizer': 'SGD'}  
-0.220775 (0.226018) with: {'optimizer': 'RMSprop'}  
-0.310961 (0.225857) with: {'optimizer': 'Adagrad'}
```

```
best_mlp_model_2 = grid_search.best_estimator_  
print(best_mlp_model_2)  
  
KerasRegressor(  
    model=<function build_model at 0x000002E011306660>  
    build_fn=None  
    warm_start=False  
    random_state=None  
    optimizer=RMSprop  
    loss=None  
    metrics=None  
    batch_size=32  
    validation_batch_size=None  
    verbose=1  
    callbacks=None  
    validation_split=0.2  
    shuffle=True  
    run_eagerly=False  
    epochs=20  
)
```


Implementing a Multilayer Perceptron

41

Fit the best model:

```
best_mlp_model_2.fit(X_train, y_train, epochs = 20, validation_data = (X_test, y_test), verbose = 1)
```

Epoch 1/20

125/125 [=====] - 2s 6ms/step - loss: 0.0601 - mae: 0.0601 - mse: 0.0064 - val_loss: 0.0341 - val_mae: 0.0341 - val_mse: 0.0018

Epoch 2/20

125/125 [=====] - 0s 4ms/step - loss: 0.0400 - mae: 0.0400 - mse: 0.0025 - val_loss: 0.0339 - val_mae: 0.0339 - val_mse: 0.0018

Epoch 3/20

125/125 [=====] - 1s 4ms/step - loss: 0.0400 - mae: 0.0400 - mse: 0.0025 - val_loss: 0.0497 - val_mae: 0.0497 - val_mse: 0.0036

Epoch 4/20

125/125 [=====] - 0s 4ms/step - loss: 0.0365 - mae: 0.0365 - mse: 0.0020 - val_loss: 0.0339 - val_mae: 0.0339 - val_mse: 0.0018

Epoch 5/20

125/125 [=====] - 0s 4ms/step - loss: 0.0354 - mae: 0.0354 - mse: 0.0019 - val_loss: 0.0346 - val_mae: 0.0346 - val_mse: 0.0019

Epoch 6/20

125/125 [=====] - 1s 4ms/step - loss: 0.0357 - mae: 0.0357 - mse: 0.0020 - val_loss: 0.0335 - val_mae: 0.0335 - val_mse: 0.0017

Epoch 7/20

125/125 [=====] - 0s 4ms/step - loss: 0.0371 - mae: 0.0371 - mse: 0.0021 - val_loss: 0.0334 - val_mae: 0.0334 - val_mse: 0.0017

Epoch 8/20

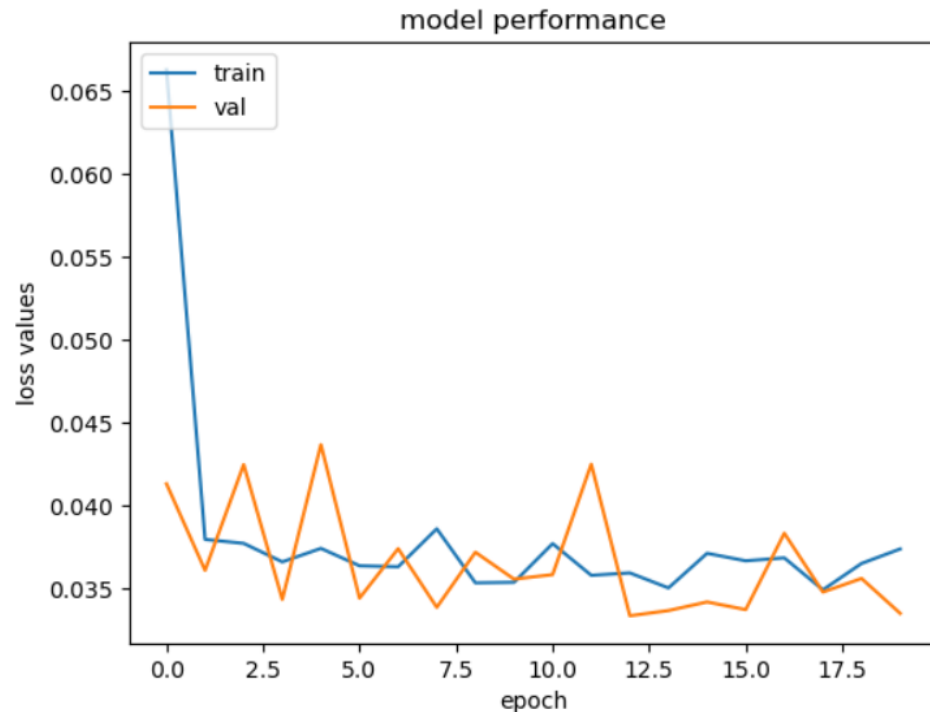
...

Implementing a Multilayer Perceptron

42

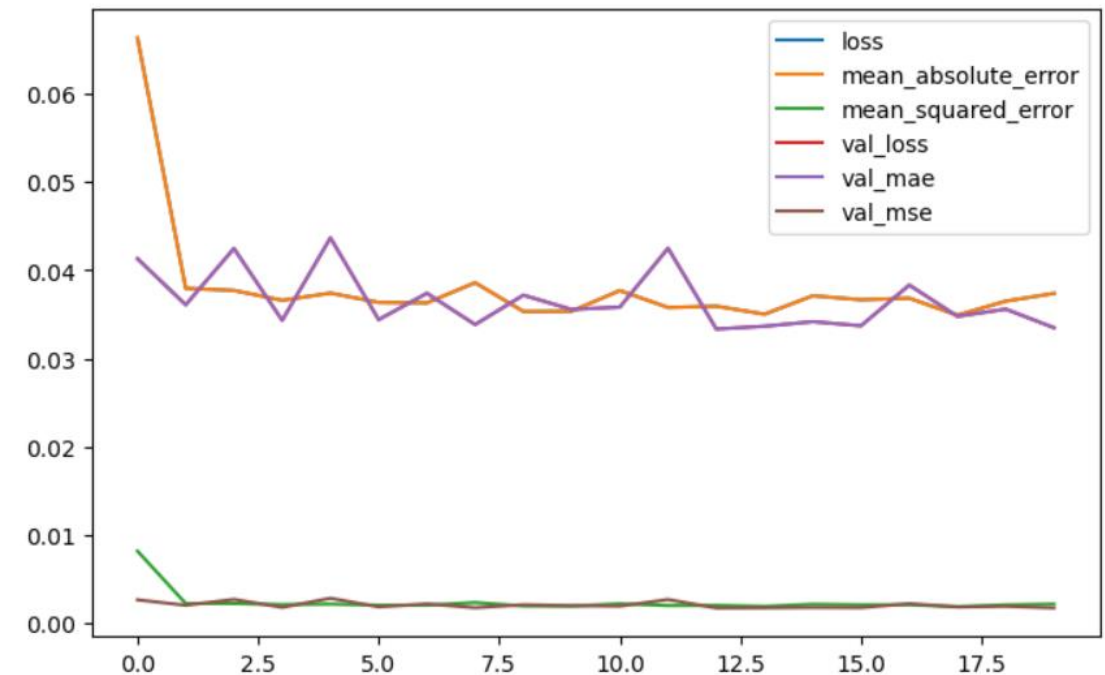
Plot the model performance:

```
plt.plot(best_mlp_model_2.history_['loss'])
plt.plot(best_mlp_model_2.history_['val_loss'])
plt.title('model performance')
plt.ylabel('loss values')
plt.xlabel('epoch')
plt.legend(['train', 'val'], loc='upper left')
plt.show()
```



Did the **model overfit**?

```
pd.DataFrame(best_mlp_model_2.history_).plot(figsize = (8, 5))
plt.show()
```



Implementing a Multilayer Perceptron

43

Predictions

Obtain the predictions:

```
predictions = best_mlp_model_2.predict(X_test)
```

```
32/32 [=====] - 0s 2ms/step
```

```
predictions = predictions.reshape(predictions.shape[0], 1)
```

Print the first five:

```
predictions[:5]
```

```
array([[0.5972543 ],  
       [0.33833626],  
       [0.26719126],  
       [0.63041043],  
       [0.29771176]], dtype=float32)
```

Predictions are scaled. We can use the `inverse_transform()` function to obtain the real values.

Unscale the predictions to see the real prices:

```
predictions_unscaled = scaler_y.inverse_transform(predictions)
```

Print the first five:

```
predictions_unscaled[:5]
```

```
array([[1481079.1 ],  
       [ 845920.5 ],  
       [ 671392.75],  
       [1562415.4 ],  
       [ 746263.44]], dtype=float32)
```

Unscale `y_test` to get the original values:

```
y_test_unscaled = scaler_y.inverse_transform(y_test)
```

Print the first five:

```
y_test_unscaled[:5]
```

```
array([[1409892.08977612],  
       [ 889385.90158426],  
       [ 635429.23051901],  
       [1613414.23305073],  
       [ 774491.65328819]])
```

Not that bad!!



Implementing a Multilayer Perceptron

44

Evaluate the model

Assess by MAE, MSE and RMSE:

```
print('MAE:', metrics.mean_absolute_error(y_test_unscaled, predictions))
print('MSE:', metrics.mean_squared_error(y_test_unscaled, predictions))
print('RMSE:', np.sqrt(metrics.mean_squared_error(y_test_unscaled, predictions)))
```


MAE: 1235057.259318577

MSE: 1637743709840.513

RMSE: 1279743.610978587

Scatter the real values with the predictions:

```
plt.scatter(y_test_unscaled, predictions)
```




Implementing a Multilayer Perceptron

45

```
def real_predicted_viz(limit):  
    plt.figure(figsize = (14, 6))  
    plt.plot(y_test_unscaled[:limit], color = 'green', label = 'Actual')  
    plt.plot(predictions_unscaled[:limit], color = 'red', label = 'Predicted')  
    plt.grid(alpha = 0.3)  
    plt.xlabel('Houses')  
    plt.ylabel('Price')  
    plt.title('Real vs Predicted')  
    plt.legend()  
    plt.show()
```

```
real_predicted_viz(200)
```



Implementing a Multilayer Perceptron



Compare the results with the ones obtained with the **Linear Regression model** created in class 4.

Which model performed better?

Hands On

47

Spyder (Python 3.6)

File Edit Search Source Run Debug Consoles Projects Tools View Help

Editor - C:\data\PythonWorkspace\dev\meanshift_algorithm.py

```
37 class Mean_Shift:
38     def __init__(self, radius=None, radius_normalize_step = 150):
39         self.radius = radius
40         self.radius_normalize_step = radius_normalize_step
41
42     def fit(self, data):
43
44         if self.radius == None:
45             all_data_centroid = np.average(data, axis=0)
46             all_data_norm = np.linalg.norm(all_data_centroid)
47             self.radius = all_data_norm/self.radius_normalize_step
48
49         centroids = {}
50
51         #initialize centroids
52         for i in range(len(data)):
53             centroids[i] = data[i]
54
55         weights = [1 for i in range(self.radius_normalize_step)]
56
57         while True:
58             new_centroids = []
59             for i in centroids:
60                 in_range = []
61                 centroid = centroids[i]
62
63                 for featureset in data:
64                     distance = np.linalg.norm(featureset-centroid)
65                     if distance == 0:
66                         distance = 0.0000000001
67                     weight_index = int(distance/self.radius)
68                     if weight_index > self.radius_normalize_step-1:
69                         weight_index = self.radius_normalize_step-1
70                     to_add = (weights[weight_index]**2)*[featureset]
71                     in_range += to_add
72
73             new_centroid = np.average(in_range, axis=0)
```

Variable explorer

Name	Type	Size	Value
batch_size	int	1	100
mnist	contrib.learn.python.learn.datasets.base.Datasets	3	Datasets object of...
n_classes	int	1	10
n_nodes_h1	int	1	500
n_nodes_h2	int	1	500
n_nodes_h3	int	1	500

Variable explorer File explorer Help

IPython console

Console 1/A

See 'tf.nn.softmax_cross_entropy_with_logits_v2'.

Epoch 0 completed out of 10 loss: 1666037.4677734375
Epoch 1 completed out of 10 loss: 377809.3128890991
Epoch 2 completed out of 10 loss: 201302.4857263565
Epoch 3 completed out of 10 loss: 119427.91378033161
Epoch 4 completed out of 10 loss: 72651.25679710507
Epoch 5 completed out of 10 loss: 45327.621502393486
Epoch 6 completed out of 10 loss: 31955.17812934518
Epoch 7 completed out of 10 loss: 23664.35610633137
Epoch 8 completed out of 10 loss: 18248.740643078025
Epoch 9 completed out of 10 loss: 19962.00065876091
Accuracy: 0.9511

In [2]:

IPython console History log

HANDS ON